Inline-style:

`![alt text](https://github.com/adam-p/markdown-here/raw/master/src/common/images/icon48.png)`

Inline-style:

Kmeans → hard assignment → ✓ GMM

Assigning Points to their own respective cluster → E

Recentering the clusters

→ Hierarchical clustering → bottom up or top-down

→ ① Centroid ② Single link ③ Complete link

④ Average link

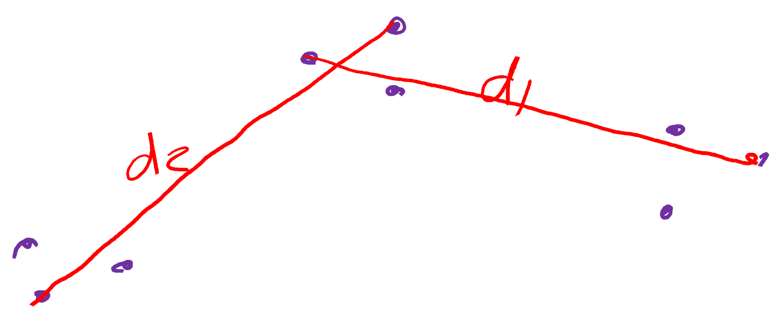

$d_1 < d_2$

Lecture 08
Density-Based Clustering

Mahdi Roozbahani
Georgia Tech

These slides are based on slides from Jing Gao, Chao Zhang
and Jiawei Han.

when you're a constant and you see d/dx

Outline

- Overview
- Basic Concepts
- The DBSCAN Algorithm
- Analysis of DBSCAN

Density-Based Clustering

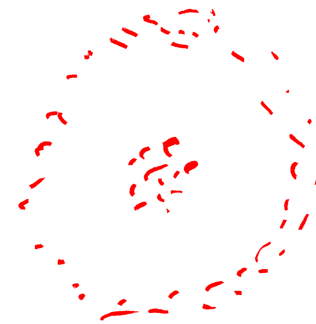

- Basic Idea
  - Clusters are dense regions in the data space, separated by regions of lower density
  - A cluster is defined as a maximal set of density-connected points
  - Detect arbitrarily shaped clusters

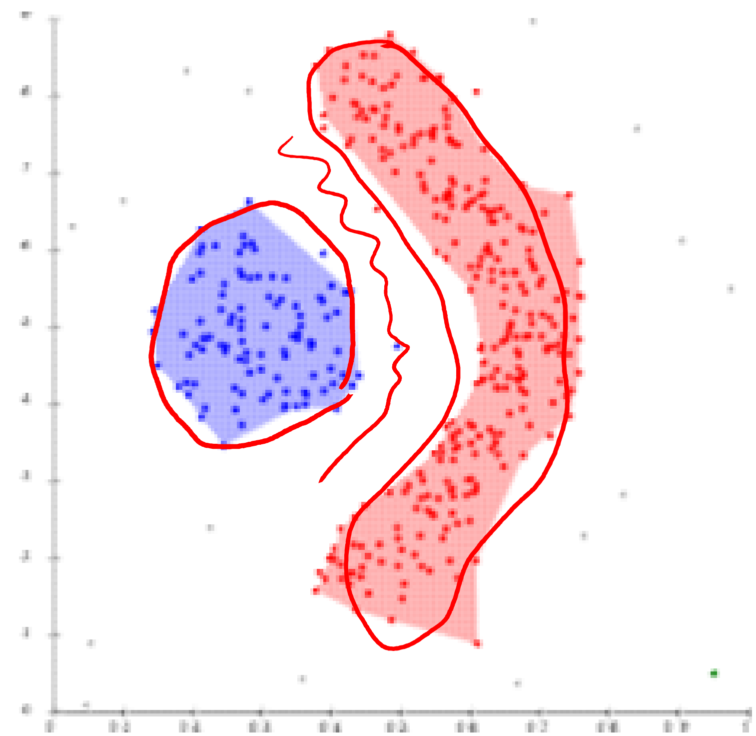

- Method
  - DBSCAN (Density-Based Spatial Clustering of Applications with Noise)

Outline

- Overview
- Basic Concepts
- The DBSCAN Algorithm
- Analysis of DBSCAN

High Density v.s. Low Density

- Two parameters
  - **Eps (ε)**: Maximum radius of the neighborhood
  - **MinPts**: Minimum number of points in the Eps-neighborhood of a point
- High density: ε-Neighborhood of an object contains at least MinPts of objects

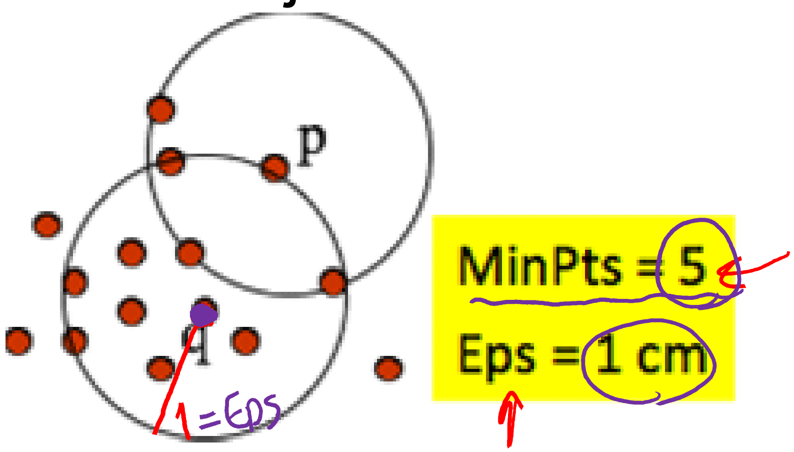

Density of ***p*** is low
Density of ***q*** is high

Core Points, Border Points, and Outliers

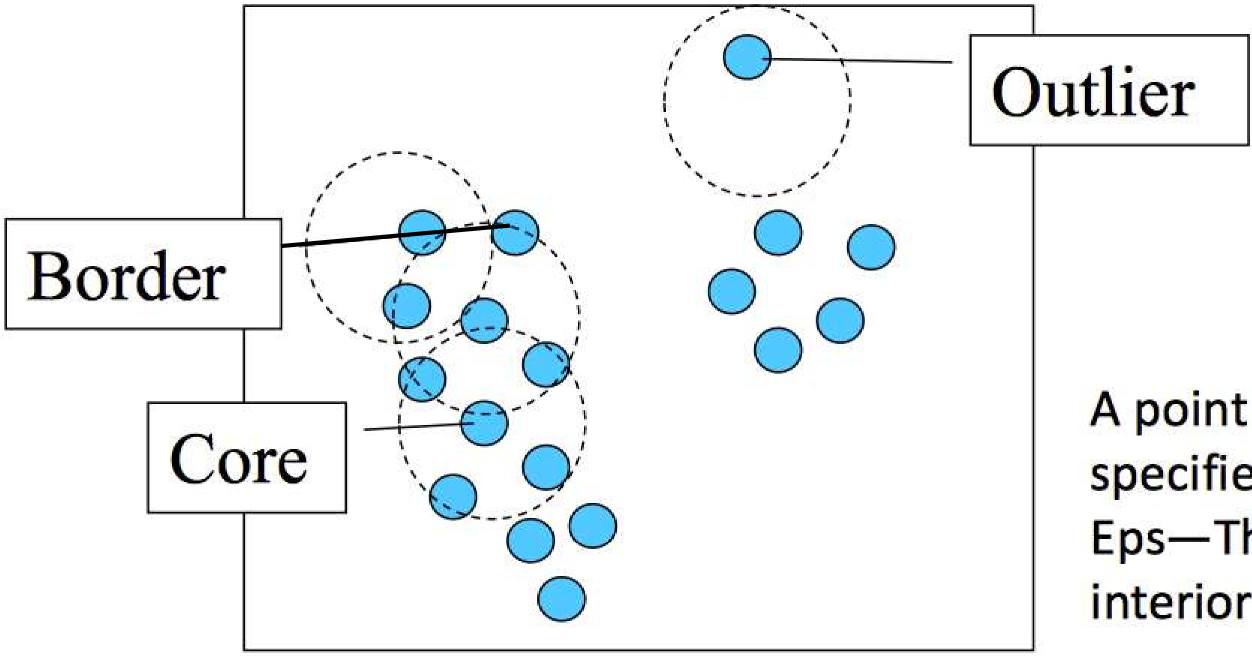

ε = 1unit, MinPts = 5

Given ε and *MinPts*, categorize the objects into three exclusive groups.

A point is a core point if it has more than a specified number of points (MinPts) within Eps—These are points that are at the interior of a cluster.

A border point has fewer than MinPts within Eps, but is in the neighborhood of a core point.

A noise point is any point that is not a core point nor a border point.

MinPts = 5

Eps = 1

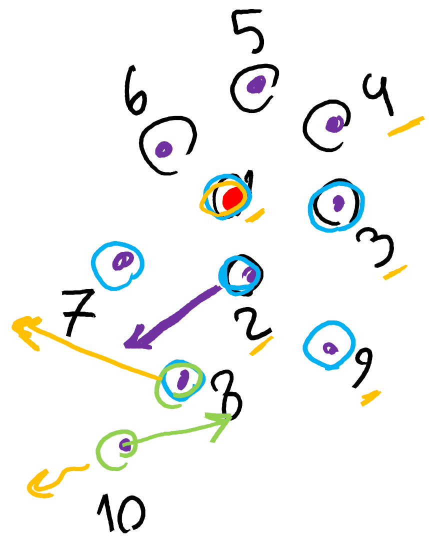

1: as a core point

$d(1,7) > Eps = 1$ X

2 is a core point

8 a border point

8 is a core point

• outlier

11

Examples

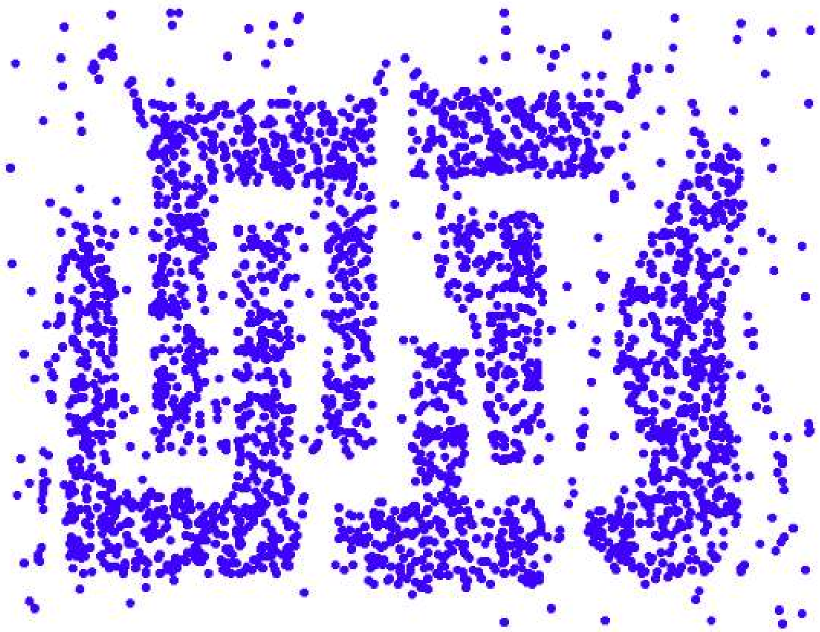

Original Points

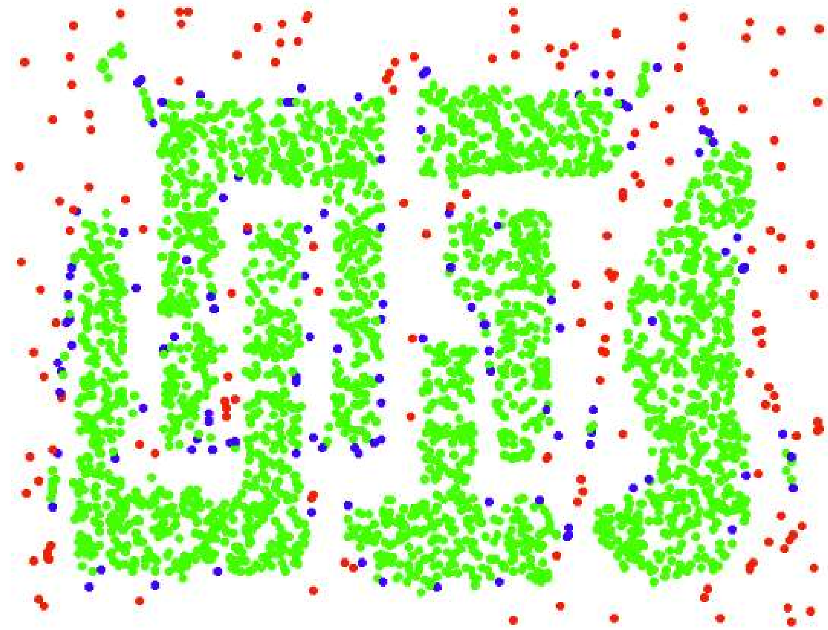

Point types: core, border and outliers

ε = 10, MinPts = 4

Density-based related points

- Direct density reachability:
  - An object p is directly density-reachable from object q if **(1) q is a core object**; and **(2) p is in q's ε-neighborhood**

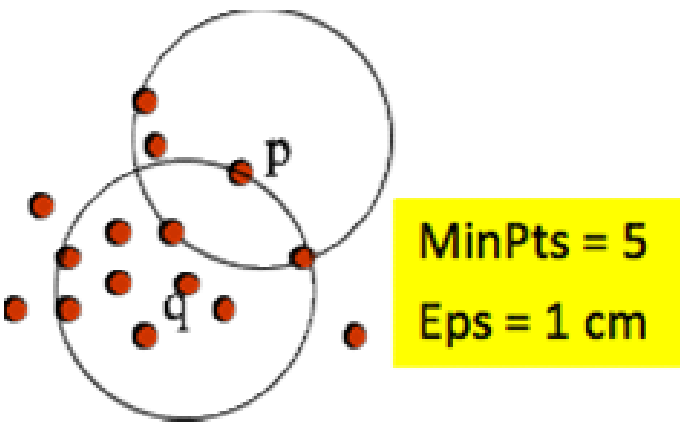

Directly Density-Reachable

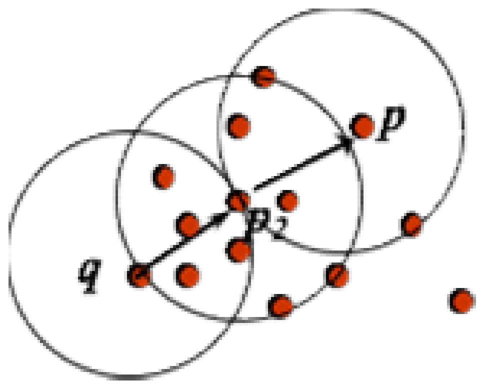

Density-Reachable

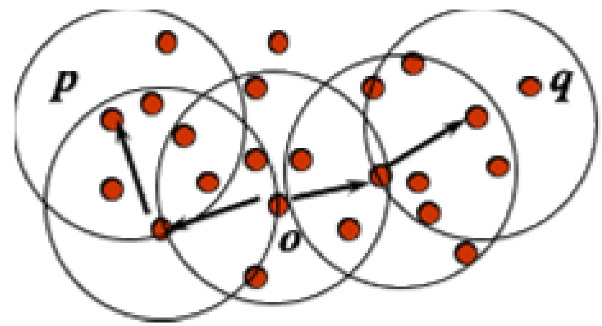

Density-Connected

Density-based related points

- Density reachability:
  - A point ***p*** is density-reachable from a point ***q*** if there is a chain of points $\boldsymbol{p_1, \ldots, p_n}$, $\boldsymbol{p_1 = q}$, $\boldsymbol{p_n = p}$ such that $\boldsymbol{p_{i+1}}$ is directly density-reachable from $\boldsymbol{p_i}$
  - $\boldsymbol{p_1 = q \rightarrow p_2 \rightarrow \ldots \rightarrow p_n = q}$

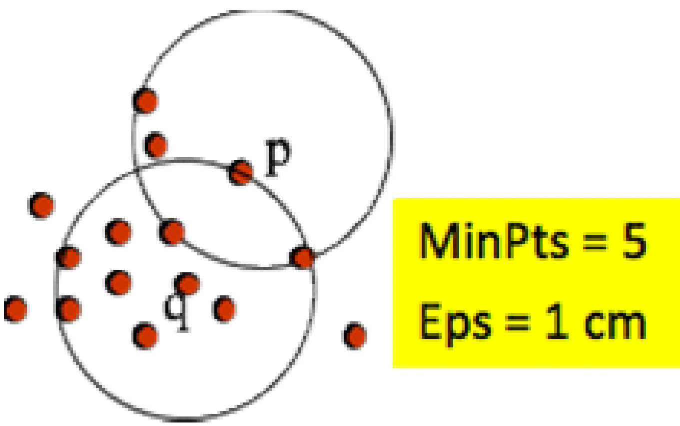

Directly Density-Reachable

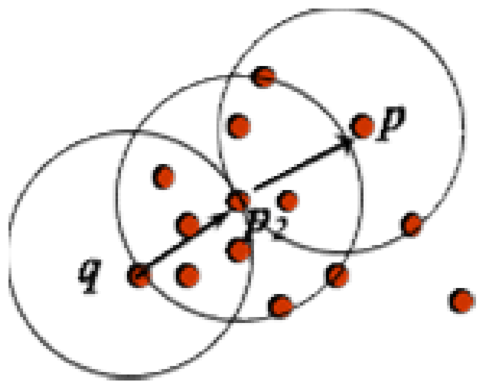

Density-Reachable

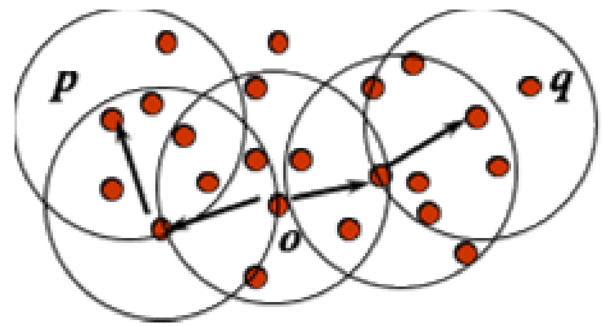

Density-Connected

Density-based related points

- Density connectivity:
  - A point ***p*** is density-connected to a point ***q*** if there is a point ***o*** such that both ***p*** and ***q*** are density-reachable from ***o***

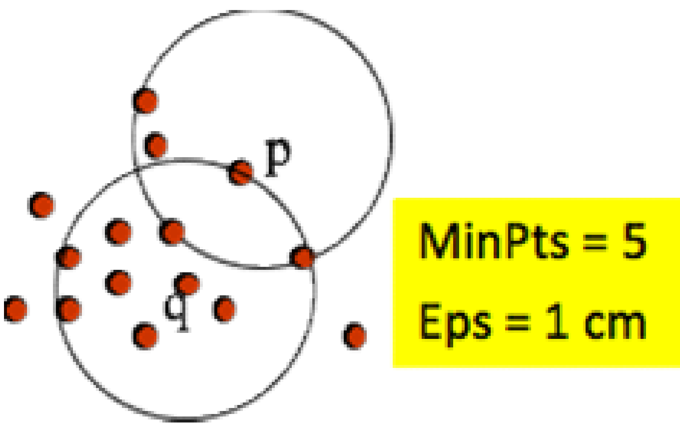

Directly Density-Reachable

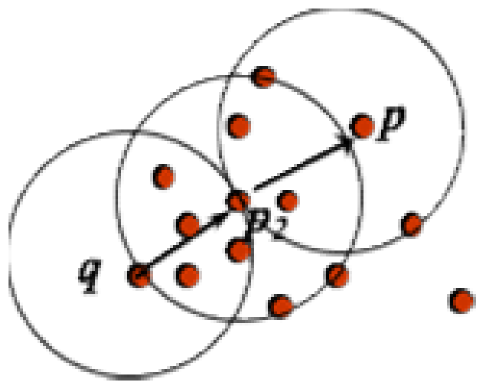

Density-Reachable

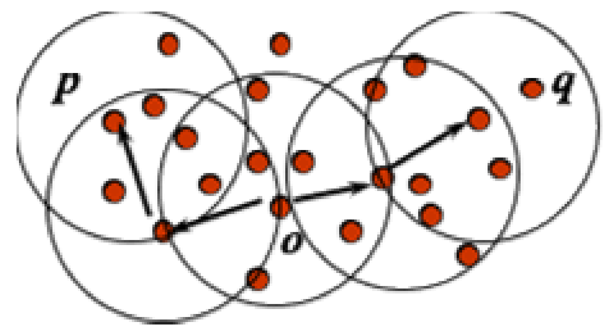

Density-Connected

Outline

- Overview
- Basic Concepts
- The DBSCAN Algorithm
- Analysis of DBSCAN

The DBSCAN Algorithm

```
DBSCAN(D, eps, MinPts)
C = 0
for each unvisited point P in dataset D
    mark P as visited
    NeighborPts = regionQuery(P, eps)
    if sizeof(NeighborPts) < MinPts
        mark P as NOISE
    else
        C = next cluster
        expandCluster(P, NeighborPts, C, eps, MinPts)

expandCluster(P, NeighborPts, C, eps, MinPts)
    add P to cluster C
    for each point P' in NeighborPts
        if P' is not visited
            mark P' as visited
            NeighborPts' = regionQuery(P', eps)
            if sizeof(NeighborPts') >= MinPts
                NeighborPts = NeighborPts joined with NeighborPts‘
        if P' is not yet member of any cluster
            add P' to cluster C

regionQuery(P, eps) return all points within P's eps-neighborhood (including P)
```

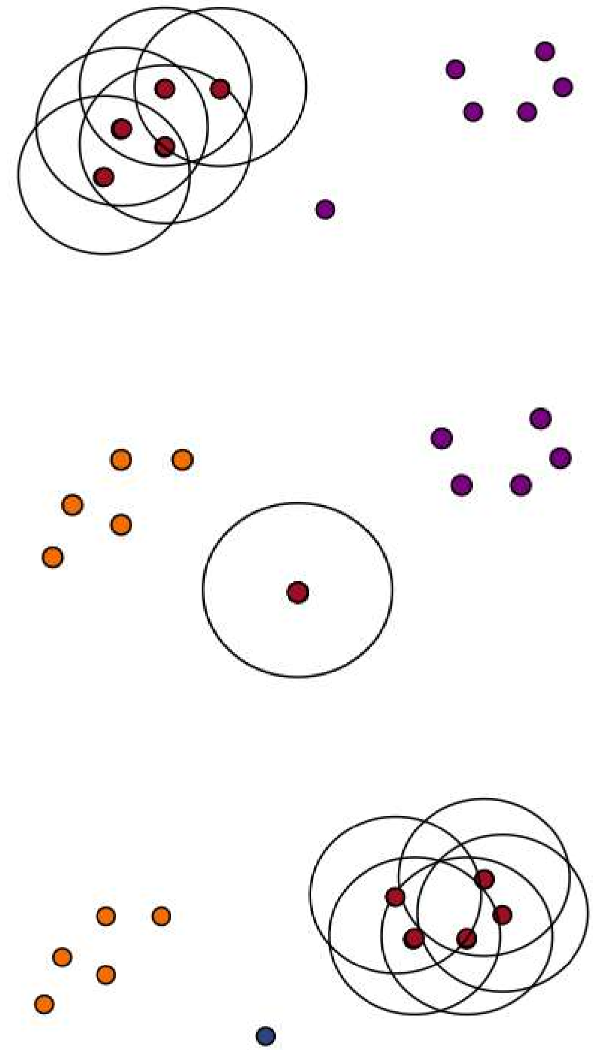

https://www.naftaliharris.com/blog/visualizing-dbscan-clustering/

Outline

- Overview
- Basic Concepts
- The DBSCAN Algorithm
- Analysis of DBSCAN

DBSCAN is Sensitive to Parameters

Figure 8. DBScan results for DS1 with MinPts at 4 and Eps at (a) 0.5 and (b) 0.4.

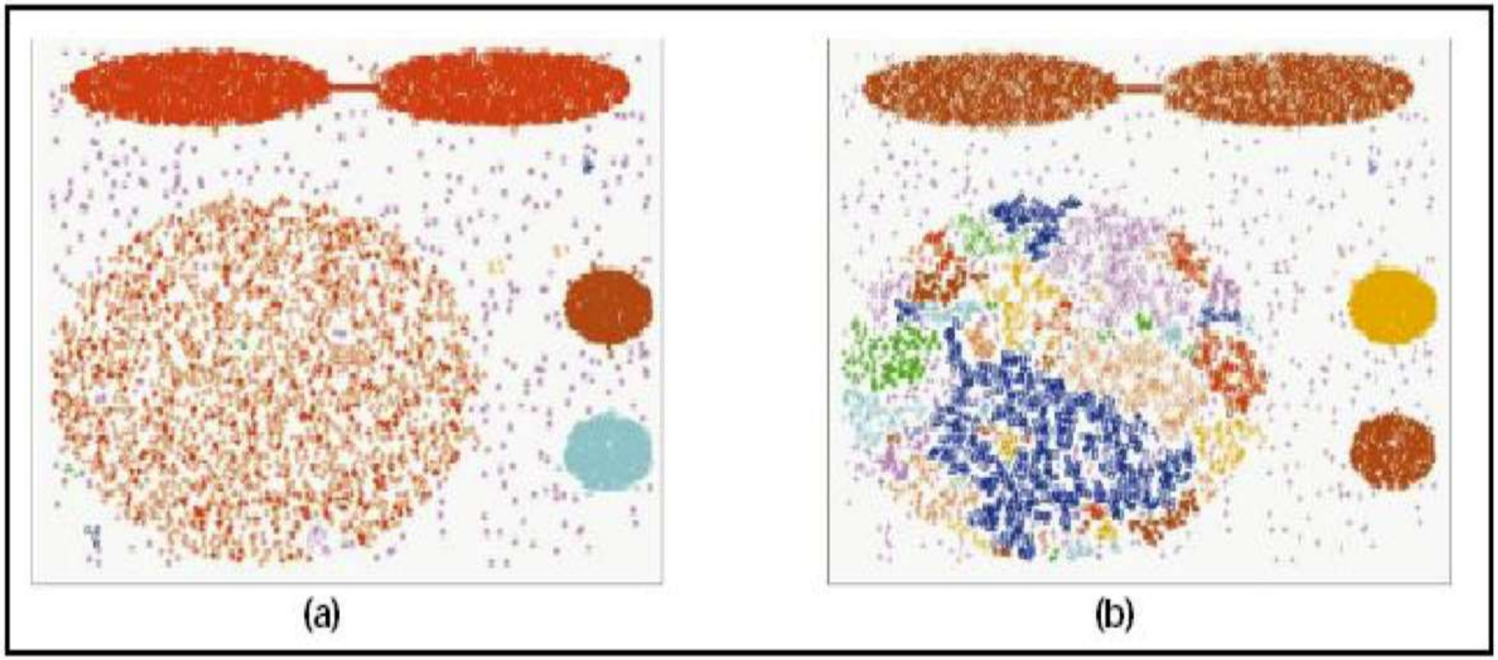

Figure 9. DBScan results for DS2 with MinPts at 4 and Eps at (a) 5.0, (b) 3.5, and (c) 3.0.

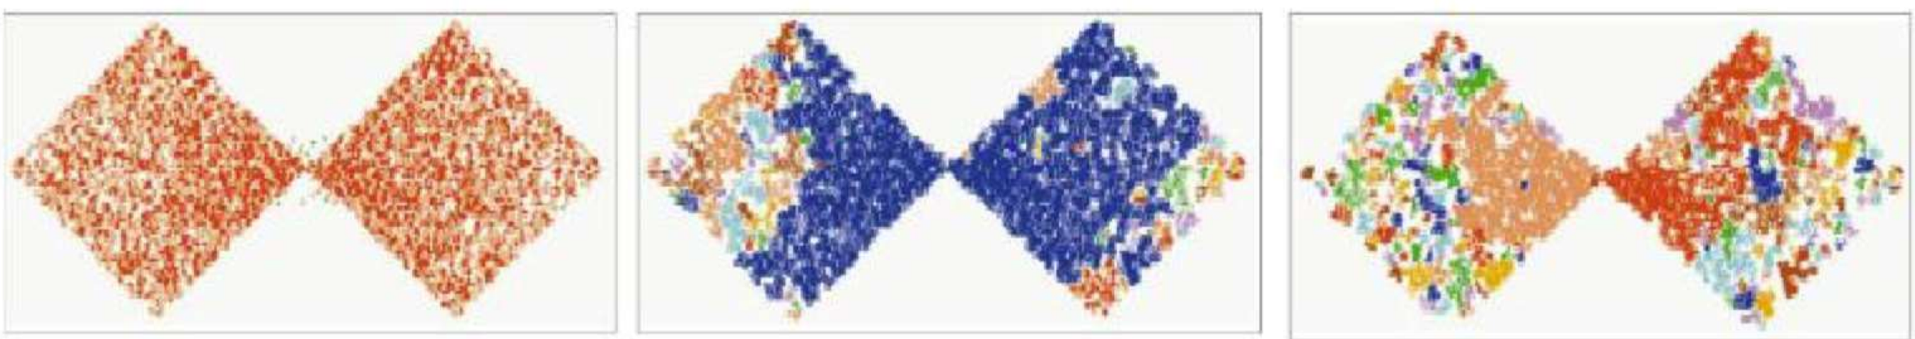

(a) (b) (c)

Image Credit: George Karypis.

ε

High value (what will happen?)

Clusters will merge and the majority of data points will be in the same cluster

Low value (what will happen?)

A large part of data won’t be clustered and considered as outliers. Because, they won’t satisfy the number of pints to create a dense region

Do we need to define the number of clusters in DBSCAN?

Nope

Minimum number of Points (**MinPts**)

MinPts = 1?

Every point will be a cluster on its own, Why?

Don't forget, in DBSCAN, a core point is counted as the number of neighboring points

MinPts = 2?

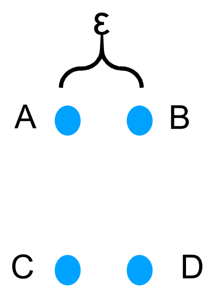

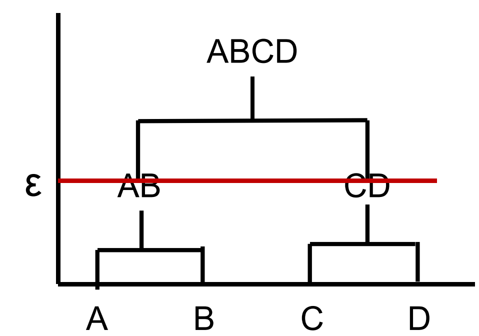

Dendogram cut at height ε

So, MinPts should be at least 3

Rule of thumb, MinPts >= D+1;
For noisy data => MinPts = 2*D (yield more significant clusters)

How about Eps? (Elbow effect)

- Idea is that for points in a cluster, their k^{th} nearest neighbors are at roughly the same distance
- Noise points have the k^{th} nearest neighbor at farther distance
- So, plot sorted distance of every point to its k^{th} nearest neighbor

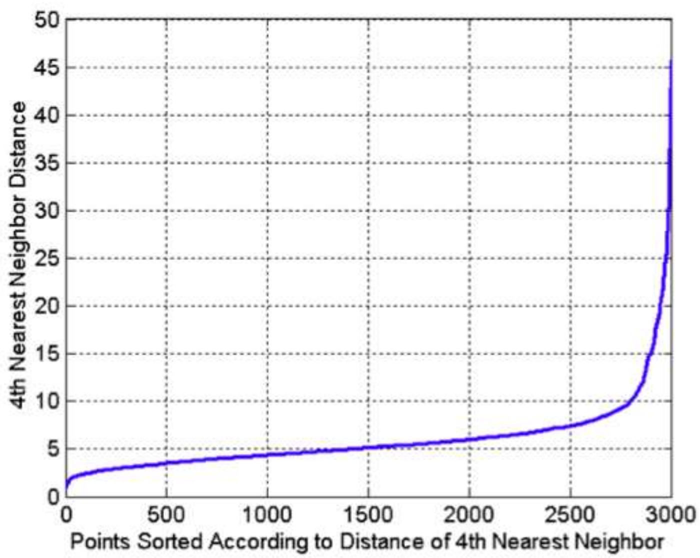

Here we have 3000 points and x-axis shows just a point index.

Point indices are sorted in ascending order based on their 4th nearest neighbor distance

Elbow effect another example

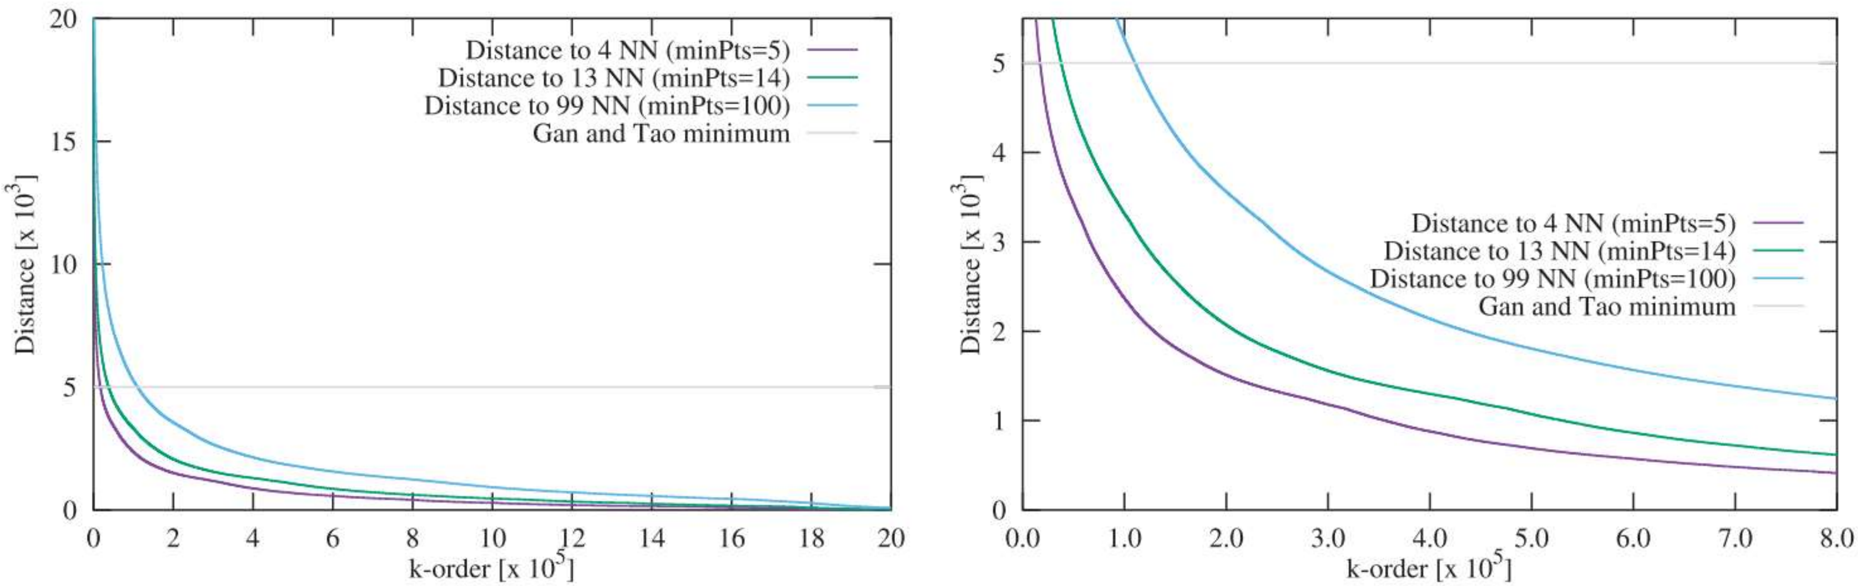

(a) k-distance plots   (b) k-distance plots (magnified region)

minPts often does Not have a significant impact on the clustering results

Erich Schuber et al

When DBSCAN Works Well

- Robust to noise
- Can detect arbitrarily-shaped clusters

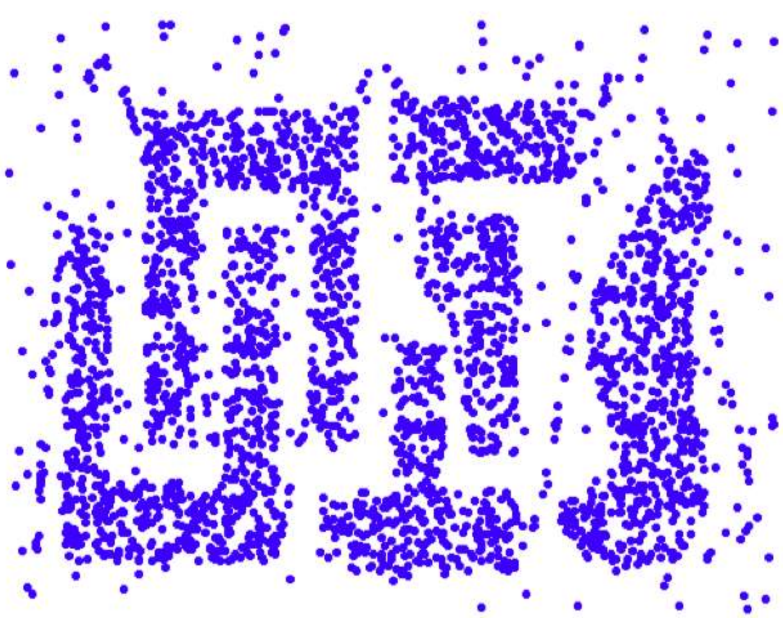

Original Points

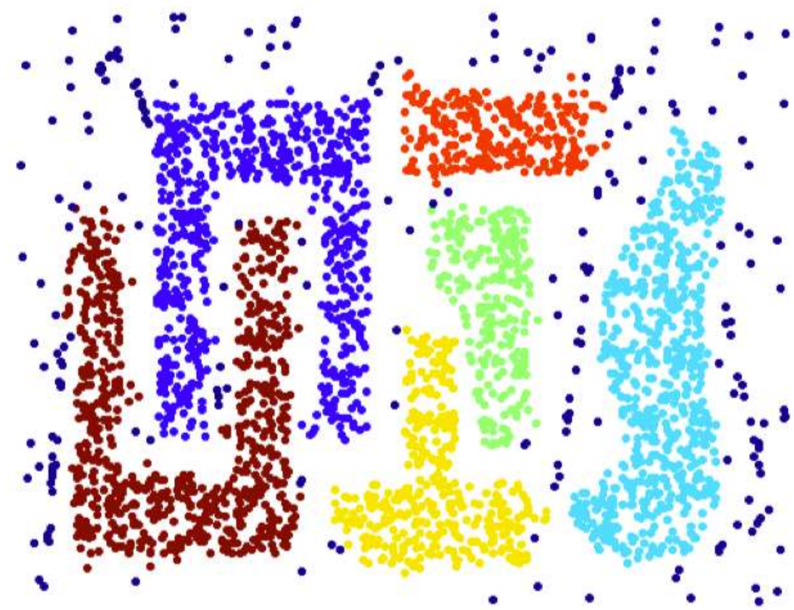

Clusters

When DBSCAN Does NOT Work Well

- Cannot handle varying densities
- Sensitive to ~~noise~~ (parameters) hard to determine the best setting of parameters

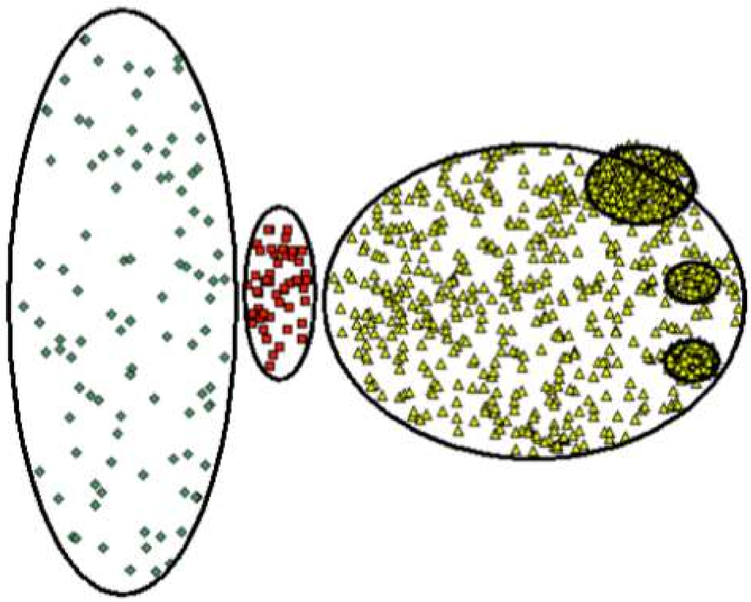

Original Points

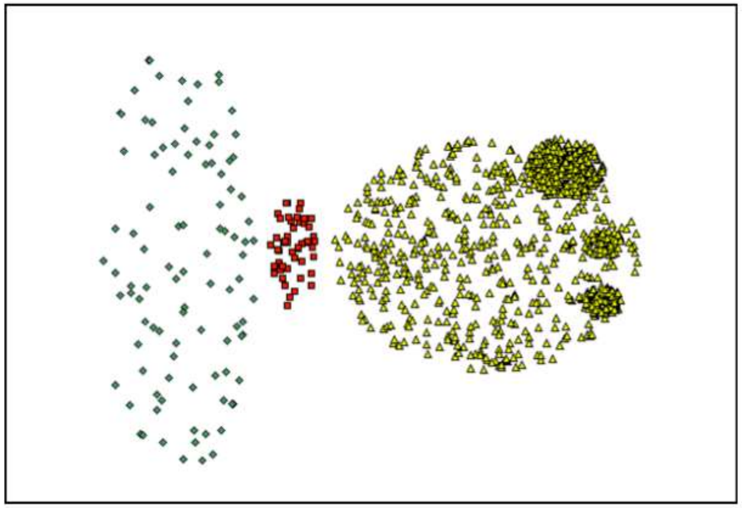

(MinPts=4, Eps=9.92).

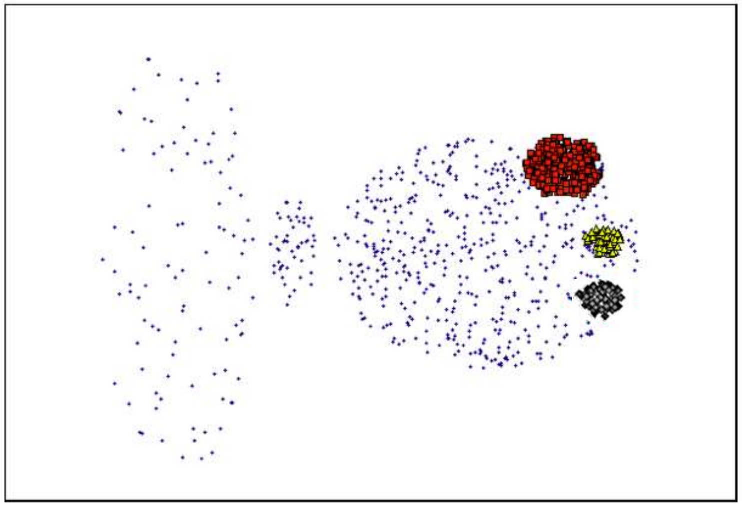

(MinPts=4, Eps=9.75)

Take-Home Messages

- The basic idea of density-based clustering
- The two important parameters and the definitions of neighborhood and density in DBSCAN
- Core, border and outlier points
- DBSCAN algorithm
- DBSCAN's pros and cons